

Efficient Synchronization of Linux Memory Regions over a Network: A Comparative Study and Implementation (Notes)

A user-friendly approach to application-agnostic state synchronization

Felicitas Pojtinger (Stuttgart Media University)

2023-08-04

Rough Structure

Rough Structure

- Abstract: A comparative analysis and implementation of various methods for synchronizing Linux memory options over a network
- Introduction
 - Examining Linux's memory management and relevant APIs
 - Use cases for memory region synchronization
- Option 1: Handling page faults in userspace with userfaultfd
 - Introduction to userfaultfd
 - Implementing userfaultfd handlers and registration in Go
 - Transferring sockets between processes
 - Examples of handler and registration interfaces (byte slice, file, S3 object)
 - Performance assessment of this approach
- Option 2: Utilizing mmap for change notifications
 - Concept: mmap a memory region with `MMAP_SHARED` to track changes in a file
 - Method 1 for detecting file changes: `inotify`
 - Limitations: `mmap` does not generate `WRITE` events

Sections/Research Questions/Ideas Brainstorming

Sections/Research Questions/Ideas Brainstorming

- Usecases: Direct Mount vs. Managed Mount vs. Migration
- Effects of high latency on different pull methods (esp. direct vs. managed)
- Effects of slow local disks or RAM on pull methods
- The asynchronous background push method (for mounts); how chunks are marked as dirty when they are being written to before the download has finished completely
- Mount backend API vs. seeder API
- Preemptive pulls and parallelized startups (n MB saved)
- Background pulling system and interface (rwat), % of availability
- Chunking system/non-aligned reads and writes, checking for correct chunking behavior
- Local vs. remote chunking
- Backend implementations, performance and usecases: File, memory, directory, dudirekta, gRPC, fRPC, Redis, S3, Cassandra

Alternative Outline

Alternative Outline

1. Abstract

- A comparative analysis and implementation of various methods for synchronizing Linux memory options over a network

2. Introduction

- 2.1 Background: Examining Linux's memory management and relevant APIs
- 2.2 Purpose: Use cases for memory region synchronization (Direct Mount, Managed Mount, Migration)
 - Discuss potential effects of high latency, slow local disks or RAM on different pull methods

3. Methodologies

- 3.1 Option 1: Handling page faults in userspace with userfaultfd
 - 3.1.1 Explanation of userfaultfd and its implementation
 - 3.1.2 Description of userfaultfd handlers and registration in Go
 - 3.1.3 The process of transferring sockets between processes
 - 3.1.4 Examples of handler and registration interfaces
 - 3.1.5 Performance assessment of this approach, with focus on pull methods and effects of system constraints

Story



- Introduction: How does memory in Linux work? Paging, swap etc.
- An introduction to mmap and how we can use it to map a file into memory/a byte slice, the role of msync
- Implementing push-based memory sync by tracking changes to a mmaped slice with polled hashing of individual chunks, why we can't use inotify, and the CPU-bound limitations of this approach
- Implementing pull-based memory sync with userfaultfd; and implementation and throughput limitations
- Implementing push-pull based memory sync with a FUSE; limitations and complexity (citing STFS)
- Implementing push-pull based memory sync with NBD; implementation of go-nbd
- Using NBD directly as a mount-based sync system with the direct mount API; limitations with latency etc., and improvements with background pulls and pushes, different backends etc., the mount

Revised Structure

1. Introduction
2. Synchronization Strategies
3. Case Studies
4. Conclusion