

---

Bachelorarbeit im Studiengang Medieninformatik

# Efficient Synchronization of Linux Memory Regions over a Network (Presentation Notes)

vorgelegt von **Felicitas Pojtinger**

an der **Hochschule der Medien Stuttgart**

am **03.08.2023**

zur Erlangung des akademischen Grades eines **Bachelor of Science**

Erstprüfer: **Prof. Dr. Martin Goik**

Zweitprüfer: **M.Sc. Philip Betzler**

---

Bachelor's Thesis

# Efficient Synchronization of Linux Memory Regions over a Network (Presentation Notes)

Author: **Felicitas Pojtinger**

University: **Hochschule der Medien Stuttgart**

Course of Study: **Media Informatics**

Date: **2023-08-29**

Academic Degree: **Bachelor of Science**

Primary Supervisor: **Prof. Dr. Martin Goik**

Secondary Supervisor: **M.Sc. Philip Betzler**

## Contents

## List of Figures

## List of Acronyms

**API** Application Programming Interface

**I/O** Input/Output

**OS** Operating System

**CPU** Central Processing Unit

**RAM** Random Access Memory

**SSD** Solid State Drive

**HDD** Hard Disk Drive

**CXL** Compute Express Link

**VFS** Virtual File System

**UUID** Universally Unique Identifier

**CRC32** Cyclic Redundancy Check 32-Bit

**LRU** Least Recently Used

**WAN** Wide Area Network

**LAN** Local Area Network

**TCP** Transmission Control Protocol

**UDP** User Datagram Protocol

**P2P** Peer-To-Peer

**NATs** Network Address Translators

**IPC** Inter-Process Communication

**RTT** Round-Trip Time

**SRP** SCSI RDMA Protocol

**GNU** GNU's Not Unix

**UNIX** UNIX Family of Operating Systems

**macOS** Apple Macintosh Operating System

**FreeBSD** Free Berkeley Software Distribution

**NBD** Network Block Device

**S3fs** S3 File System

**NVMe** Non-Volatile Memory Express

**LTFS** Linear Tape File System

**LTO** Linear Tape-Open

**EXT4** Fourth Extended Filesystem

**Btrfs** B-Tree File System

**LTFS** Linear Tape File System

**ELF** Executable and Linkable Format

**C** C Programming Language

**Rust** Rust Programming Language

**Go** Go Programming Language

**C++** C++ Programming Language

**ARM** ARM RISC Computer Processor Architecture

**x86** x86 CISC Computer Processor Architecture

**RISC-V** RISC-V RISC Computer Processor Architecture

**LPDDR5** Low-Power Double Data Rate 5

**HTTP** Hypertext Transfer Protocol

**HTTPS** HTTP Secure

**HTTP/2** HTTP Version 2

**QUIC** Quick UDP Internet Connections

**WebRTC** Web Real-Time Communication

**Wasm** WebAssembly

**WASI** WebAssembly System Interface

**IETF** Internet Engineering Task Force

**OIDC** OpenID Connect

**AWS** Amazon Web Services

**CNCF** Cloud Native Computing Foundation

**S3** Simple Storage Service

**TLS** Transport Layer Security

**mTLS** Mutual TLS

**SSH** Secure Shell

**DoS** Denial of Service

**JSON** JavaScript Object Notation

**JSONL** JSON Lines

**SQL** Structured Query Language

**NoSQL** Not Only SQL

**Protobuf** Protocol Buffers

**IDL** Interface Definition Language

**DSL** Domain-Specific Language

**KV** Key-Value

**Syscalls** System Calls

**VM** Virtual Machine

**RPC** Remote Procedure Call

**REST** Representational State Transfer

**FUSE** File Systems in Userspace

- Introduction
  - Title slide
  - ToC
  - About me
  - Abstract/introduction
- Methods
  - Pull-based synchronization with `userfaultfd`/Userfaults in Go with `userfaultfd`
    - \* Technology section: Memory organization & hierarchy
    - \* Technology section: Page faults
  - Push-based synchronization with `mmap` and hashing/file-based synchronization, discussion
    - \* Technology section: ‘mmap’
    - \* Technology section: Delta synchronization

- Push-based synchronization with FUSE/FUSE implementation in Go, discussion
  - \* Technology section: FUSE
- Mounts with NBD/NBD with go-nbd
  - \* Technology section: NBD
- Push-Pull Synchronization with Mounts/managed mounts with r3map
  - \* Technology section: RTT, LAN and WAN
- Pull-Based Synchronization with Migrations/Live migration
  - \* Technology section: Pre- and post-copy VM migration, workload analysis
- Optimizations
  - Pluggable Encryption, Authentication and Transport
  - Concurrent Backends
  - Remote Stores as Backends
  - Concurrent RPC frameworks (dudirekta) and connection pooling (gRPC)
- Discussion and Results
  - Testing Environment
  - Access methods (userfaults vs. direct vs. managed mounts): Latency & Throughput, discussion
  - Initialization: Polling vs. udev
  - Chunking methods: Local vs. remote
  - RPC frameworks; discussion
  - Backends: Latency & throughput; discussion
  - General limitations of the r3map library (deadlocks etc.)
- Implemented Use Cases
  - Using mounts for remote swap with [ram-dl](#)
  - Mapping tape into memory with tapisk
- Future Use Cases
  - Improving cloud storage clients
  - Universal database, media and asset streaming
  - Universal app state mounts and migrations
- Conclusion
- Thanks